Summaries: quality versus latency

Qwen3.5-9B precision series and separate reference configurations

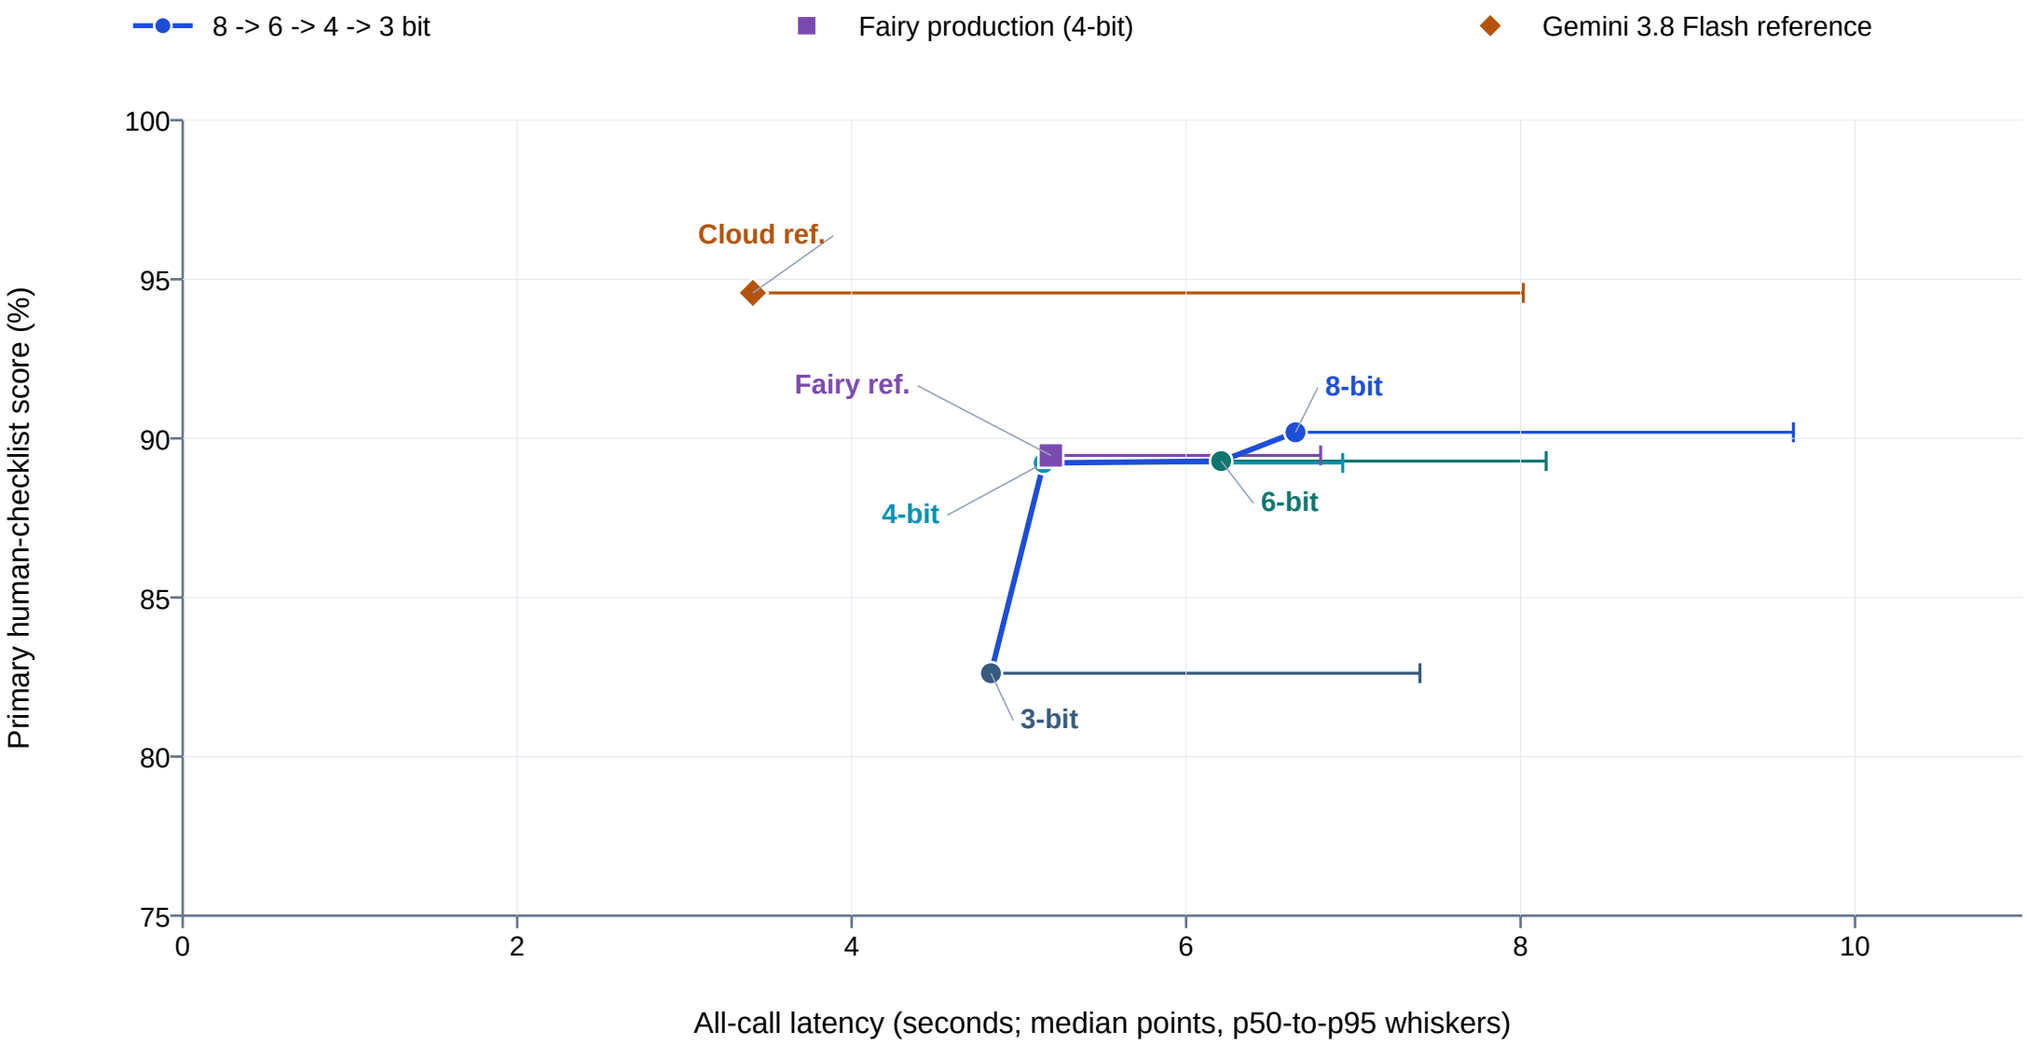

10 invented-text summaries x 5 repeats = 50 measured calls per configuration.

One blinded human checklist review; original max_words included. Sensitivity scores are separate.

Quality averages repeats within each prompt, then weights prompts equally; repeats are not independent prompts.

Whiskers show latency spread, not confidence intervals. Quality axis is zoomed to 75-100%.

All completed calls are included, including swapping. References are outside the registered knee sequence.

Local thinking: off. Cloud thinking: low; successful-attempt latency excludes retries and backoff.